

A Polls-of-Polls Forecast for the 2024 United States Presidential Election Through Generalized Linear Model*

The Impact of Pollster Differences, Population Type, and Poll Recency on Predicting Support for Kamala Harris

Jiwon Choi Kevin Roe

November 4, 2024

The 2024 U.S. Presidential Election has far-reaching global implications for economic, environmental, and social policy, making it a focal point for international attention. This paper applies a Bayesian generalized linear model to analyze trends in Kamala Harris’ support, accounting for pollster differences, population types, and poll recency. Despite steady support at 48%, Harris is projected to lose the national popular vote and key battleground states to Donald Trump. Our findings emphasize the significance of polling methodologies and their role in predicting U.S. policy outcomes. Future research should incorporate the Electoral College to address current limitations.

Table of contents

1	Introduction	2
2	Data	3
2.1	Measurement	3
2.2	Outcome variables	4
2.2.1	The Predicted Proportion of Support a Candidate Received in a Poll . .	4
2.3	Predictor variables	6
2.3.1	Type of Pollster	6
2.3.2	Population Type	7
2.3.3	Poll Recency	9

*Code and data are available at: https://github.com/jwonc4602/2024_US_Election_Forecast.

3	Model	11
4	Results	12
5	Discussion	13
5.1	How has the Forecast Changed Over Time?	13
5.2	Which States Prefer Harris?	14
5.3	Implications	16
5.4	Weaknesses and Next Steps	17
A	Appendix	19
A.1	YouGov Methodology Analysis	19
A.2	Idealized Methodology	19
A.3	Diagnostics for model	20
	References	22

1 Introduction

The United States Presidential Election is one of the most consequential events of 2024. As a key country in international relations and the largest economy in the world, the results of the U.S. election determine the country’s domestic and foreign policy for the next four years, which will have a significant impact on important initiatives such as tackling climate change or international conflicts. In anticipation of the 2024 U.S. election, this paper aims to predict the possible outcomes of the election by analyzing the level of support that Kamala Harris will gain.

We forecast support for Kamala Harris based on polling results and apply a Bayesian generalized linear model. The main parameter of interest is the proportion of votes or support that Harris received in surveys, which is traced over time. By considering the results from different poll-making organizations and other demographic factors, our objective is to correct for variation across different voter bases.

Factoring in the different results of poll-making organizations, our model found that specific pollsters or demographic characteristics increase or decrease Harris’ predicted voter share. Analyzing the data at a national and state level, we predict that Harris has a minority in the national popular vote and key swing states, leading us to conclude that Donald Trump will win the 2024 US Presidential Election. Understanding the trajectory the US might make in foreign policy, economic, or global politics prediction allows global stakeholders to take preemptive measures for policy changes of the newly elected government. Thus, this study is a strong tool to navigate uncertainty around the 2024 election.

The paper is structured as follows: Section 2 and Section 3 explores the data and methodology used, highlighting the filtering and modeling techniques applied to the data; Section 4 presents

the results from the Bayesian generalized linear model; Section 5 discusses the broader implications of our findings; and Section A highlights YouGov’s methodology, outlines an idealized survey methodology, and shows model diagnostics.

2 Data

We have used a poll-of-polls of the upcoming US presidential election dataset obtained from FiveThirtyEight (FiveThirtyEight, n.d.). Data was collected and analyzed using R statistical programming software (R Core Team 2023), with additional packages like tidyverse (Wickham et al. 2019), rstanarm (Goodrich et al. 2022) knitr (Xie 2020), here (Müller 2020), and many others for support. The dataset includes a wide range of poll results from various polls, with key variables such as pollster, sample size, percentage of support for Harris, and the date conducted for polls.

In performing the analysis, we used several R packages. We used tidyverse (Wickham et al. 2019) for data manipulation and visualization, and rstanarm (Goodrich et al. 2022) for Bayesian modeling. To visualize results, we used ggplot2 (Wickham 2016) and kableExtra (Zhu 2020) to format tables for presentation.

2.1 Measurement

FiveThirtyEight aggregates various poll results from national and state-wide polls, showing 17765 observations (FiveThirtyEight, n.d.) as of November 3, 2024. The poll takes a sample of the electorate and asks for the voters’ candidate of choice. By factoring polls from a state and national level, FiveThirtyEight aims to represent the public perception of the two candidates during the election to predict the election’s outcome.

Each poll aims to predict an actual event. The raw data is susceptible because each pollster has varying polling methods. There are other limitations such as sampling error, response error, or distortion from participants misunderstanding the question.

Beyond filtering for high-quality pollsters, we filtered for various criteria to prepare the analysis data. From the raw data, we filtered out any missing numbers, cleaned names, and removed all polls other than Harris, effectively generating prediction support for Harris. Our dataset contains qualified polls for both national and state-level polls. Regardless, selection bias and sampling bias are still a concern because polls only represent a part of the population. Finally, polling methodology differences can introduce bias in the study. Overall, we use the data from individual responses to election polls and filter the data for analysis to predict who will win the popular vote based on data from the pollsters.

2.2 Outcome variables

2.2.1 The Predicted Proportion of Support a Candidate Received in a Poll

The main variable we aim to forecast is the `pct` variable, which represents the proportion of the vote a candidate received in a poll. Table 1 and Figure 1 show the summary statistics and distribution of the `pct` variable in a filtered dataset that only comprises votes from relatively high-quality polling organizations. We also get the popular vote predictions for each candidate using the `predict()` function and add this to the dataset as the `predicted_pct` variable. We have shown the summary statistics and distribution for `predicted_pct` in Table 2 and Figure 2, respectively. Comparing our predictions to the data, our predictions have a smaller range than the cleaned data. The difference in the mean is 2% lower for our prediction than our cleaned data. The results suggest there is less variation in our predictions than in the cleaned dataset.

Table 1: Summary statistics for the proportion of support candidates received in the poll

mean	median	min	max	sd	n
45.25	46	18	70	4.82	7268




Figure 1: Distribution of the proportion of support candidates received in the poll

Table 2: Summary statistics for the predicted proportion of support candidates received in the poll

mean	median	min	max	sd	n
47.07	46.83	45.67	50.44	1.39	311

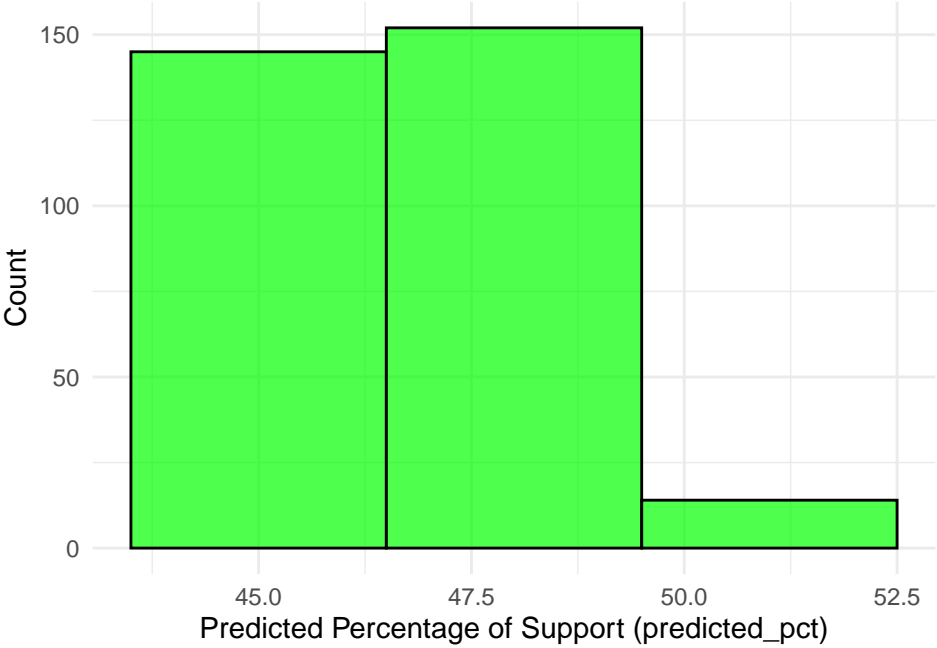


Figure 2: Distribution of the predicted proportion of support candidates received in the poll

2.3 Predictor variables

2.3.1 Type of Pollster

The pollster variable was selected to consider the effect of changes in the pollster on the observed variable. The pollster variable represents the polling organization that conducted the poll. The distribution of polling counts for different pollsters in Figure 3 suggests that the data is dominated by two pollsters: YouGov and Siena/NYT. Further analysis is needed in their polling methodology to determine potential biases.

Table 3: Number of unique high-quality polling organizations

pollster	Count
CES / YouGov	28
Marquette Law School	24
McCourtney Institute/YouGov	1
Siena/NYT	136
The Washington Post	10
YouGov	110
YouGov Blue	1
YouGov/Center for Working Class Politics	1




Figure 3: Distribution of high-quality polling organizations

2.3.2 Population Type

Population Type was distinguished among different groups surveyed into two groups surveyed: ‘Voters’ and ‘Adults’. ‘Adults’ is a more general poll that includes voters, likely voters, and non-voters. Figure 4 shows that the majority of the polls surveyed voters rather than the general adult population. The discrepancy between the number of polls that surveyed voters and adults could introduce potential biases in responses, as those who are voters may have stronger opinions than those who are not.

Table 4: Number of unique high-quality polling organizations

population	Count
Adults	22
Voters	289




Figure 4: Distribution of population type

2.3.3 Poll Recency

Poll Recency, shown through the variable `recent_poll`, labels any polls conducted in the last 30 days as Recent, and any polls collected before that are labeled Older. Based on our results shown in Table 5, 19 more polls have been collected in the last 30 days than before. This introduces biases in responses as recent responses are more represented in our model. In rapidly changing political environments, public opinion can shift dramatically due to major events, debates, or crises, making recent data crucial for accurate modeling. The distribution of this variable can also be found in Figure 5.

Table 5: Number of polls that were collected within 30 days and after

<code>recent_poll</code>	Count
Older	146
Recent	165

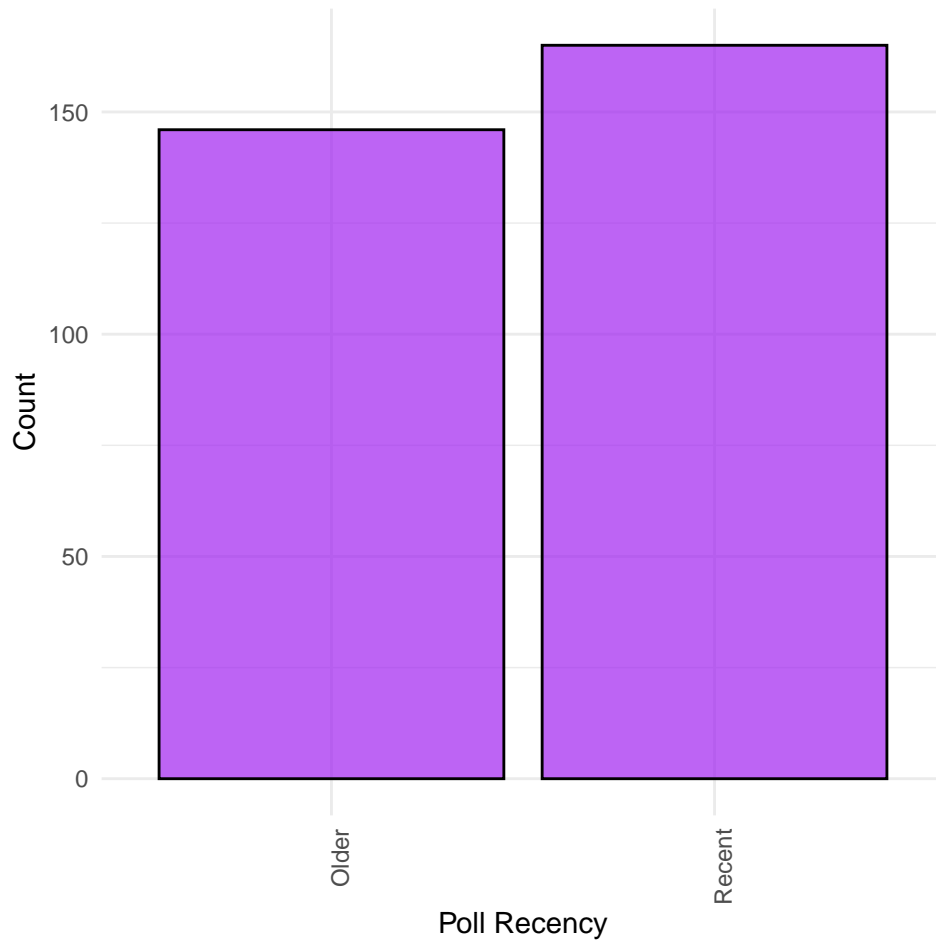


Figure 5: Distribution of on the amount of polls collected in 30 days

3 Model

For our analysis, we employ a Bayesian generalized linear model (GLM) to forecast the popular vote percentage for Kamala Harris. This approach allows us to capture polling characteristics and account for known variations between pollsters, population types, and the recency of polling data. By incorporating these factors, we aim to model an estimate of support for Harris.

The first step in our process involved selecting a reliable dataset for model development. Here, we utilized high-quality national polling data gathered after Harris’s campaign announcement. We filtered the dataset to include polls with a numeric grade of 3.0, ensuring data quality, and focused only on polls conducted after July 21, 2024, the date of Harris’s declaration of candidacy. The filter helped limit bias from older polls from when Joe Biden was the Democratic nominee, which may not reflect the current voter sentiment.

The GLM is specified as follows:

$$Y_i = \beta_0 + \beta_1 x_{pollster_i} + \beta_2 x_{population_i} + \beta_3 x_{recentpoll_i} + \epsilon_i \quad (1)$$

In equation 1, each β represents a coefficient determined through regression analysis. The variables chosen for this project are pollster, population type, and recency of the poll. Each predictor variable was carefully selected based on its significance in polling analysis and its correlation with voting trends. The identity of the polling organization is important to our model, as each pollster may exhibit unique biases. Including Pollster as a fixed effect allows us to account for these variations without introducing unnecessary complexity. Population Type distinguishes among different groups surveyed (e.g., voters, likely voters) and helps in capturing generalizability. Categorizing polls as either recent or older ensures that more recent polls are better predictors of voting behavior closer to the election. Y_i denotes the predicted popular vote percentage for Kamala Harris in the i -th poll. ϵ_i is the Gaussian-distributed error term, accounting for residual variation in the model.

To enhance the model, Bayesian priors were applied, introducing regularization and incorporating plausible ranges grounded. For the coefficient priors β , a normal distribution with a mean of 0 and a scale of 2.5 (autoscaled) was chosen to provide flexibility while mitigating overfitting. Similarly, the intercept uses a normal prior with a mean of 0 and a scale of 2.5 to stabilize model estimates. For the error term (sigma), an exponential prior with a rate of 1 was selected to constrain the residuals, aligning with Gaussian assumptions.

The model was implemented in R (R Core Team 2023) using the `rstanarm` package, which offers an accessible interface for Bayesian generalized linear models (GLMs), allowing specification of priors and customization of model parameters. Once the logistic regression model for predicting the popular vote is developed, we apply the `predict()` function in R (R Core Team 2023) to generate popular vote percentage predictions for Harris. This prediction is then added to the dataset as a new variable (`predicted_pct`) for further analysis. The results are

Table 6: Coefficients from the GLM Model

term	estimate	std.error	conf.low	conf.high
(Intercept)	48.64	0.88	47.18	50.12
pollsterMarquette Law School	-1.66	1.12	-3.56	0.23
pollsterMcCourtney Institute/YouGov	-2.48	3.66	-8.44	3.56
pollsterSiena/NYT	-4.76	0.92	-6.28	-3.25
pollsterThe Washington Post	-3.67	1.42	-5.99	-1.39
pollsterYouGov	-2.56	0.90	-4.04	-1.12
pollsterYouGov Blue	-4.48	3.60	-10.54	1.26
pollsterYouGov/Center for Working Class Politics	-3.76	3.63	-9.69	2.23
populationVoters	1.87	0.95	0.33	3.41
recent_pollRecent	-0.07	0.43	-0.77	0.63

saved in a CSV file (`vote_for_Harris_predictions.csv`), enabling an examination of potential outcomes in the vote. This approach facilitates an understanding of vote distributions at a national level and provides a foundation for forecasting electoral outcomes based on polling data.

4 Results

To assess model reliability, we examined several key diagnostics. Convergence metrics, such as Rhat values, were very close to 1 for all parameters, indicating strong convergence. Additionally, the effective sample size (`n_eff`) was high across parameters, suggesting low autocorrelation and contributing to model stability. See more details of model diagnostic here: Section A.3.

The model assumes residuals follow a Gaussian distribution, though this may not fully capture extreme polling variances. Pollster effects are treated as fixed to simplify the model and reduce overfitting, even if it limits responsiveness to shifts in polling methods. A hierarchical model with random pollster intercepts was considered, but its added complexity did not enhance accuracy. The selected GLM specification balances interpretability and performance, making it the preferred choice for this study.

Table 6 presents the estimated coefficients for the predictors in our GLM model. These coefficients fit into the GLM equation, allowing us to interpret the impact of each predictor on Harris’s predicted vote percentage. Positive values indicate a higher predicted percentage, while negative values indicate a decrease. Key predictors, such as specific pollsters and population types, show distinct effects on the forecasted outcomes.

Figure 6 represents the model coefficients, with error bars indicating the confidence interval for each estimate. Positive coefficients suggest that specific pollsters or demographic characteristics increase Harris’s predicted vote share. For example, the ‘populationVoters’ shows a positive impact, while most of the pollsters have a slight negative effect. These error bars help contextualize the reliability of each predictor.

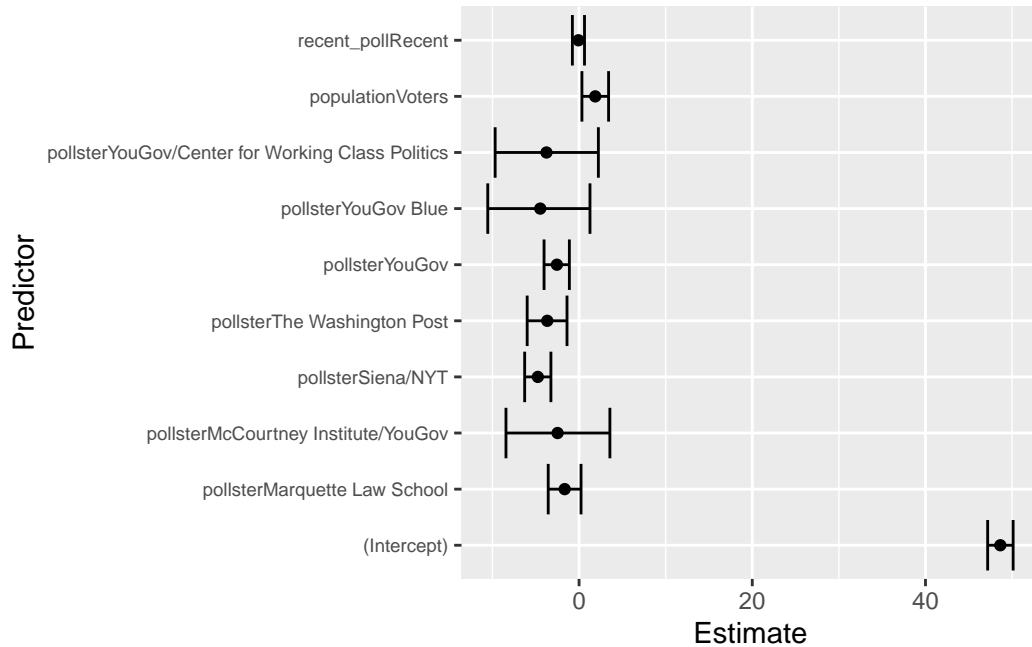


Figure 6: Coefficient Estimates for Predictors

5 Discussion

5.1 How has the Forecast Changed Over Time?

Figure 7 maps the predicted and real approval percentages for Kamala Harris during the time of her campaign. In Figure 7, our model shows that predicted approval ratings during Kamala Harris' campaign sat at around 47%, on average, which is in line with current ABC polling methods. However, if we examine the real polling approval rating percentages, it shows that Kamala Harris' approval rating across polls significantly varied over time, but remained below 50%. The observed proportion of support over time varies because different polls have different methodologies and suggest that there was significant variation in Kamala Harris' polling results during her campaign.

At the beginning of her campaign, Harris' approval rating significantly rose, but it eventually leveled out to just around 47%. Her initial rise in the polls came from excitement from the replacement of Joe Biden and the added momentum from a new candidate. Leading up to the election, Harris' approval ratings saw spikes and dips, most likely due to key catalysts such as the Presidential Debate, Hurricane Milton, and the escalation in the Ukraine-Russian War and the Israel-Palestine Conflict. However, leading up to the election, Figure 7 shows Harris' approval rating settling just approximately at 47%, signaling that Harris is likely to lose the popular vote, or at very best, the 2024 US presidential election will be close.

The most apparent observation from Figure 7 is that Harris’ approval rating is quite volatile over time. The reason for this volatility is due to the differences between pollsters’ methodologies, leading to drastically different results. Our model factors for these differences and notes the different pollsters as a notable predictor variable. The apparent variation in Figure 7 suggests the influence methodology and time have on the results of the poll. As time progressed, Harris’ approval rating also varied, signaling that any candidate’s momentum can change rapidly and that maintaining a stable base of support can be challenging in a dynamic environment.




Figure 7: Comparison of Predicted vs Actual National Poll Percentages for Harris Over Time

5.2 Which States Prefer Harris?

However, unlike other countries, the United States’ election runs through the Electoral College. To increase representation for smaller states, the Electoral College led to the formation of “swing states”, which are states that have similar levels of support among voters and are important to the outcome of the election. Notable swing states include Wisconsin, Pennsylvania, Arizona, Minnesota, Georgia, and Florida. Therefore, to have a clearer picture of the outcome of the United States election, it is important to do a state-by-state analysis to understand Harris’ approval ratings across key states.

In our model, by separating state-specific polls, we can analyze and model the poll results by states where state-specific polls were conducted. Examining Figure 8, common battleground

states such as Arizona, Pennsylvania, Wisconsin, and Michigan are polled more frequently at the state level than partisan states such as California and New York. Examining the results at these key battleground states, the results show that Harris' popularity recently fell below 50% during October and November in key states such as Arizona, Michigan, Pennsylvania, Nevada, and Wisconsin. Though Harris is predicted to have a majority in Ohio and Florida, the disappointing results in the other states make her lead in these states less powerful.

Figure 9 shows Harris with a slight lead of 50.4% in the popular vote across five key states: Virginia, Ohio, New Hampshire, Minnesota, and Florida. In other states, she is expected to either be in a close race or lose the popular vote. Pennsylvania, Nevada, Michigan, and Georgia are crucial battlegrounds, where Harris's minority share and Trump's projected lead reduce her chances of winning the election. Thus, our model predicts that Harris is unlikely to win the general popular vote or key swing states. Consequently, we anticipate Donald Trump will win the 2024 U.S. Presidential election.

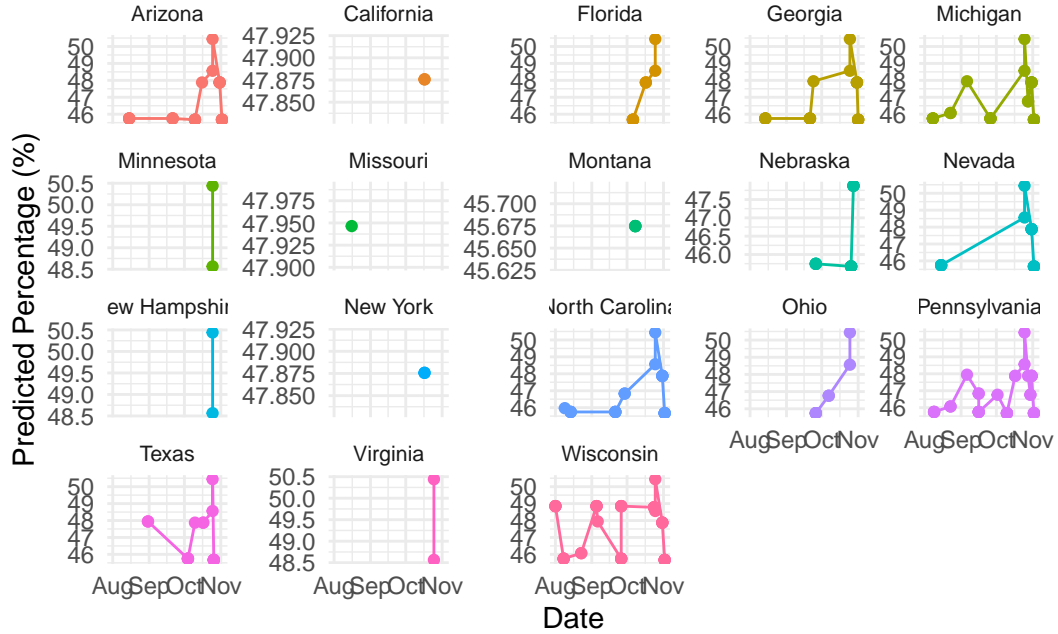


Figure 8: State-by-State comparison of predicted percentages over time

5.3 Implications

If our prediction is correct and Trump returns to office, we will see a shift in the United States' domestic and foreign policy compared to the Biden administration. Donald Trump's presidency from 2016 to 2020 was marked by tax cuts, stronger border enforcement, isolationist foreign policy, and a massive rollback on environmental regulations (BBC News 2024). During this campaign cycle, Trump has largely promoted similar policies, representing a stark contrast to Harris' campaign. The paper implies that if Trump wins the election as projected, then we will see the re-introduction of these policies and a different America on the global stage. As the United States election is one of the most important international events of this year, predicting the event's outcome is critical for foreign governments and businesses to guide strategy on US relations and regulations.

Moreover, the paper has lessons for the Harris campaign. Harris should prioritize swing states where small shifts in public opinion can determine electoral outcomes. Harris' campaign may need to adjust its messaging and outreach in battleground states to mitigate her projected losses. Also, as shown through Figure 7, popularity can fluctuate based on public enthusiasm and external factors. Campaigns must focus on mobilizing their voter base effectively in the weeks leading up to the election. Though the paper can help guide the Harris campaign leading up to the election, it also can serve to help future candidates in guiding their strategy surrounding shaping public opinion and prioritizing swing states.




Figure 9: Summary of Kamala Harris' predicted popular vote by key states

5.4 Weaknesses and Next Steps

The first weakness in our study is the discrepancy of data points between states. Figure 9 shows that certain states have more frequent state-specific polls than others. While states like California, Montana, and New York have historically leaned towards one party than the other, the discrepancy in data points between states is a weakness in our data and hurts our ability to conclude our data, as it is difficult to make out trends. Therefore, for future studies, the next potential step could be to conduct more state-specific polls and map out a candidate's popularity across each state.

The second weakness of the study is its reliance on the popular vote, which may not accurately reflect the results of the election. The United States does not elect its President based on the popular vote, but it uses the Electoral College. While the popular vote is a valuable metric, the Electoral College system means that a candidate can lose the popular vote and still win the presidency, as seen in past elections. Harris' projected loss in key states, combined with a predicted minority in the national polls, signals that simply gaining a majority in the popular vote may not be sufficient for electoral success under the current system. Therefore, the next step for this study is to further incorporate the Electoral College and its methodology into the study, requiring more state-specific polling and adjusting the estimand.

Finally, as we are working with polling data, there is potential for bias. As mentioned earlier in Section 2.3.3, the potential for bias may disproportionately influence model predictions.

This has broader implications for the accuracy of electoral forecasts and highlights the need for continually updated polling to provide a more accurate reflection of voter sentiment as Election Day approaches.

A Appendix

A.1 YouGov Methodology Analysis

YouGov utilizes a panel-based methodology for conducting online surveys and recruiting participants who voluntarily join their platform. Survey topics range widely, covering areas such as political views, policy opinions, and voter behavior. For election polling, YouGov employs a Multilevel Regression and Post-stratification (MRP) model that integrates respondent data with external sources like voter files to enhance the accuracy of outcome predictions (YouGov 2024a).

Surveys target U.S. adults, with samples carefully selected based on survey topics. The sampling process begins with YouGov’s online panel, comprised of individuals who join voluntarily. To improve representativeness, YouGov establishes demographic quotas by age, gender, and political affiliation and applies MRP-based post-survey adjustments. These steps help to mitigate demographic imbalances and minimize sampling bias (YouGov 2024a).

While YouGov’s non-probability sampling method offers a fast and cost-effective approach to data collection, it may introduce self-selection bias. Though MRP adjustments use broad population data to reduce these biases, reliance on a non-random sample can still result in skewed outcomes. Survey format and design also play roles in influencing participant responses, potentially impacting representativeness across groups with differing levels of digital literacy or social trust (YouGov 2023).

Non-response bias presents additional challenges, as individuals who abstain from participation may differ in ways from those who complete surveys. YouGov addresses this issue by applying post-survey weighting to align responses with population characteristics and by using imputation to estimate missing data based on observed patterns (YouGov 2024b). Despite these steps, non-response bias can still affect predictions, particularly when groups less inclined to participate, such as those with lower levels of education or trust, are underrepresented (CloudResearch, n.d.).

YouGov’s combination of a panel-based approach and MRP modeling provides a strong framework to manage non-random sampling challenges and mitigate non-response bias. However, limitations persist, particularly regarding self-selection and question phrasing. While well-suited for large-scale political surveys, this methodology requires careful interpretation of sampling and non-response strategies (YouGov 2024a).

A.2 Idealized Methodology

Because the U.S. operates on the Electoral College system, where state affiliation matters, we would use Stratified Random Sampling to ensure each state is represented proportionally in the survey, improving accuracy and representativeness. My target population would be

voter-eligible citizens in the United States who are above 18 years old, with a sampling frame utilizing voter registration databases stratified by key demographics, including age, race, gender, geographic region, income, and education level. We aim for a sample size of around 10,000 respondents to achieve a reasonable margin of error, ensuring findings, though we expect the actual sample may be less due to various factors.

To recruit respondents, we would primarily utilize online methods, such as surveys on social media platforms like Facebook and Instagram to engage younger audiences, and interactive voice response systems for older or less tech-savvy individuals. To ensure the data collected is valid, we will cross-check responses with voter registration databases and implement logic checks within the survey to detect inconsistencies. For instance, if a respondent claims to have already voted but indicates they're unlikely to vote, this response will be flagged as unreliable.

Addressing non-response bias is important, and we plan to over-sample underrepresented groups while employing multiple attempts to contact individuals online and conducting weekly polls. Frequent surveys will allow me to aggregate results using a moving average to smooth out short-term fluctuations in responses. we will also implement checks to prevent duplicate entries in the online survey and track phone responses to ensure unique participants. Given that the U.S. does not conduct elections based on a popular vote, we will weight the survey results to reflect the voting population as indicated by U.S. Census data and voter turnout estimates, considering demographic factors such as race, gender, and age.

To allocate the \$100,000 budget, we will spend \$20,000 on a survey platform subscription, \$50,000 on targeted digital ads and phone surveys for recruitment, \$20,000 on data analysts and poll aggregation services, and \$10,000 for re-contact surveys and miscellaneous costs. However, we recognize that this budget may be on the lower end to ensure a high-quality poll and that the actual costs might exceed this amount. Most expenses will arise from data collection, and we have accounted for this in our budget planning.

Please find the proposed question list for the online survey attached [here](#).

A.3 Diagnostics for model

Figure 10 compares observed data (dark line) with replicated posterior predictions (lighter lines). The close alignment suggests that the model accurately captures the data's central tendency and variability. Figure 11 and Figure 12 show that the sampling algorithm used, the Markov chain Monte Carlo (MCMC) algorithm, did not run into issues as the posterior distribution for the model was created. Using the checks presented by Alexander (2023), both graphs do not show anything abnormal since the trace plots in Figure 11 display substantial horizontal fluctuation across chains, indicating good mixing, while the Rhat values in Figure 12 are close to 1 and well below 1.1, further supporting convergence.




Figure 10: Posterior Predictive Check: Comparison of Observed and Replicated Data

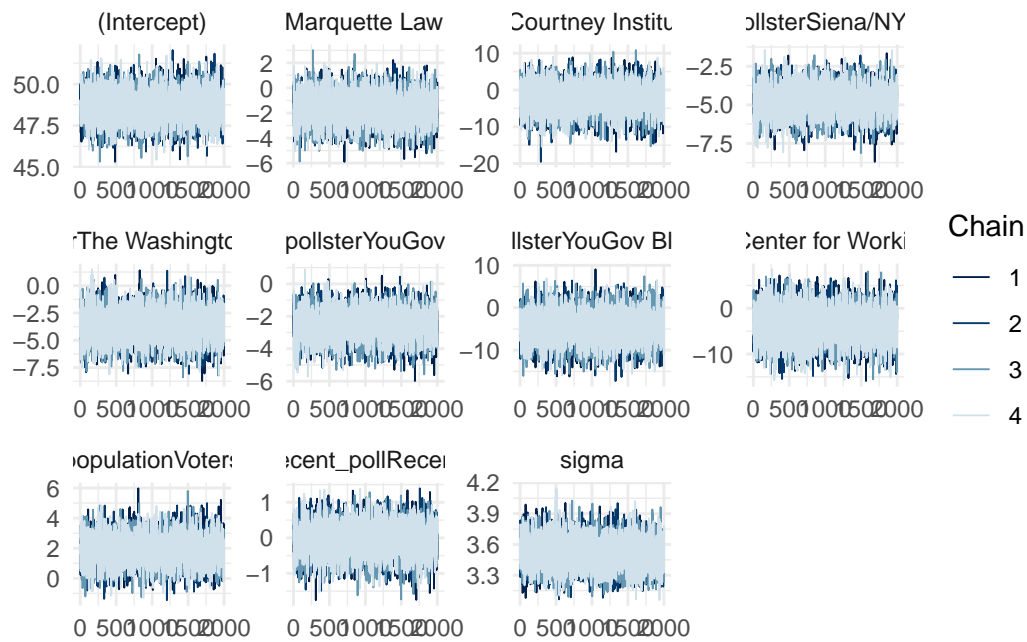


Figure 11: Checking the convergence of the MCMC algorithm - Trace

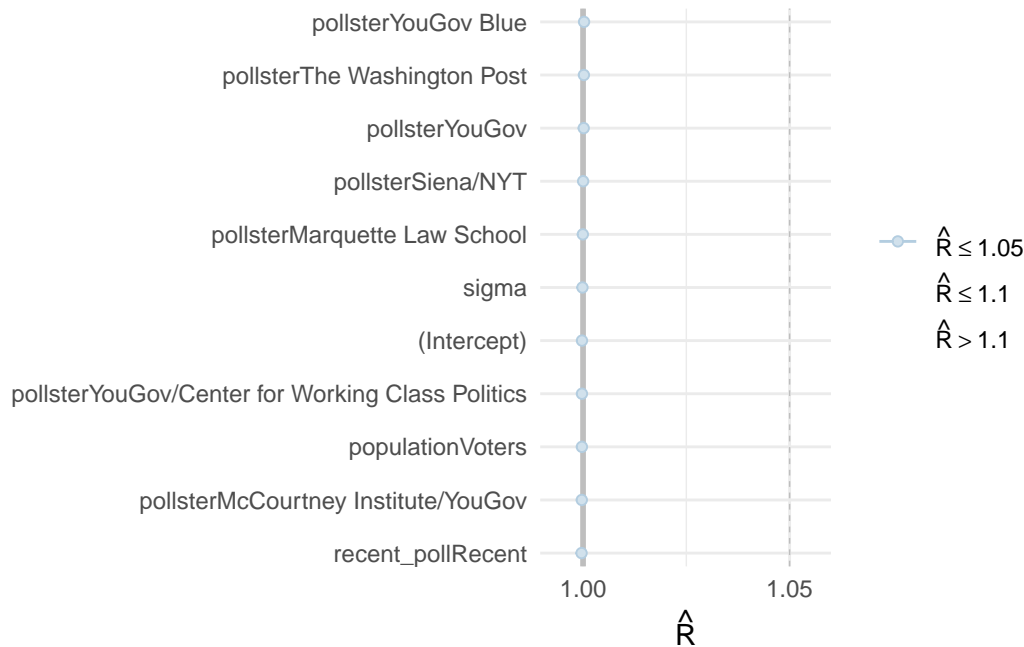


Figure 12: Checking the convergence of the MCMC algorithm - Rhat

References

- Alexander, Rohan. 2023. “Telling Stories with Data.” Chapman; Hall/CRC. <https://tellingstorieswithdata.com/>.
- BBC News. 2024. “Where Donald Trump Stands on 10 Key Issues, from Tax to Israel.” BBC News website. <https://www.bbc.com/news/articles/c2l1qd1nwnjo>.
- CloudResearch. n.d. “What Are the Effects of Non-Response Bias in Surveys?” <https://www.cloudresearch.com/resources/blog/what-effects-non-response-bias-surveys/>.
- FiveThirtyEight. n.d. “Latest Polls - FiveThirtyEight.” <https://projects.fivethirtyeight.com/polls/>.
- Goodrich, Ben, Jonah Gabry, Imad Ali, and Sam Brilleman. 2022. “Rstanarm: Bayesian Applied Regression Modeling via Stan.” <https://mc-stan.org/rstanarm/>.
- Müller, Kirill. 2020. *here: A Simpler Way to Find Your Files*. <https://CRAN.R-project.org/package=here>.
- R Core Team. 2023. *R: A Language and Environment for Statistical Computing*. Vienna, Austria: R Foundation for Statistical Computing. <https://www.R-project.org/>.
- Wickham, Hadley. 2016. *ggplot2: Elegant Graphics for Data Analysis*. Springer-Verlag New York. <https://ggplot2.tidyverse.org>.
- Wickham, Hadley, Mara Averick, Jennifer Bryan, Winston Chang, Lucy D’Agostino McGowan, Romain François, Garrett Golemund, et al. 2019. “Welcome to the tidyverse.” *Journal*

- of Open Source Software* 4 (43): 1686. <https://doi.org/10.21105/joss.01686>.
- Xie, Yihui. 2020. *Knitr: A General-Purpose Package for Dynamic Report Generation in r*. <https://yihui.org/knitr/>.
- YouGov. 2023. “How Does Poll Methodology Affect Survey Response?” <https://today.yougov.com/politics/articles/46070-poll-methodology-question-affect-survey-response>.
- . 2024a. “How YouGov’s MRP Model Works for the 2024 Presidential and Congressional Elections.” <https://today.yougov.com/politics/articles/50587-how-yougov-mrp-model-works-2024-presidential-congressional-elections-polling-methodology>.
- . 2024b. “YouGov Panel Methodology.” <https://today.yougov.com/about/panel-methodology>.
- Zhu, Hao. 2020. *kableExtra: Construct Complex Table with 'Kable' and Pipe Syntax*. <https://CRAN.R-project.org/package=kableExtra>.